

```
In [1]: import findspark  
findspark.init()
```

```
In [2]: import pyspark
```

```
In [3]: from pyspark.sql import SparkSession
```

```
In [4]: spark = SparkSession.builder.appName('BigdataProject').getOrCreate()
```

```
In [5]: spark
```

Out[5]: **SparkSession - in-memory**

SparkContext

[Spark UI](#)

Version	v3.3.1
Master	local[*]
AppName	BigdataProject

```
In [6]: df = spark.read.option('header', 'true').csv('CIS_Automotive_Kaggle_Sample.csv', inferSchema=True)
```

```
In [7]: df.columns
```

```

Out[7]: ['vin',
        'stockNum',
        'firstSeen',
        'lastSeen',
        'msrp',
        'askPrice',
        'mileage',
        'isNew',
        'color',
        'interiorColor',
        'brandName',
        'modelName',
        'dealerID',
        'vf_ABS',
        'vf_ActiveSafetySysNote',
        'vf_AdaptiveCruiseControl',
        'vf_AdaptiveDrivingBeam',
        'vf_AdaptiveHeadlights',
        'vf_AdditionalErrorText',
        'vf_AirBagLocCurtain',
        'vf_AirBagLocFront',
        'vf_AirBagLocKnee',
        'vf_AirBagLocSeatCushion',
        'vf_AirBagLocSide',
        'vf_AutoReverseSystem',
        'vf_AutomaticPedestrianAlertingSound',
        'vf_AxleConfiguration',
        'vf_Axles',
        'vf_BasePrice',
        'vf_BatteryA',
        'vf_BatteryA_to',
        'vf_BatteryCells',
        'vf_BatteryInfo',
        'vf_BatteryKWh',
        'vf_BatteryKWh_to',
        'vf_BatteryModules',
        'vf_BatteryPacks',
        'vf_BatteryType',
        'vf_BatteryV',
        'vf_BatteryV_to',
        'vf_BedLengthIN',
        'vf_BedType',
        'vf_BlindSpotMon',
        'vf_BodyCabType',
        'vf_BodyClass',
        'vf_BrakeSystemDesc',
        'vf_BrakeSystemType',
        'vf_BusFloorConfigType',
        'vf_BusLength',
        'vf_BusType',
        'vf_CAN_AACN',
        'vf_CIB',
        'vf_CashForClunkers',
        'vf_ChargerLevel',
        'vf_ChargerPowerKW',
        'vf_CoolingType',
        'vf_CurbWeightLB',
        'vf_CustomMotorcycleType',
        'vf_DaytimeRunningLight',
        'vf_DestinationMarket',
        'vf_DisplacementCC',
        'vf_DisplacementCI',
        'vf_DisplacementL',
        'vf_Doors',
        'vf_DriveType',
        'vf_DriverAssist',
        'vf_DynamicBrakeSupport',
        'vf_EDR',
        'vf_ESC',
        'vf_EVDriveUnit',
        'vf_ElectrificationLevel',
        'vf_EngineConfiguration',
        'vf_EngineCycles',
        'vf_EngineCylinders',
        'vf_EngineHP',

```

```
'vf_EngineHP_to',
'vf_EngineKW',
'vf_EngineManufacturer',
'vf_EngineModel',
'vf_EntertainmentSystem',
'vf_ForwardCollisionWarning',
'vf_FuelInjectionType',
'vf_FuelTypePrimary',
'vf_FuelTypeSecondary',
'vf_GCWR',
'vf_GCWR_to',
'vf_GVWR',
'vf_GVWR_to',
'vf_KeylessIgnition',
'vf_LaneDepartureWarning',
'vf_LaneKeepSystem',
'vf_LowerBeamHeadlampLightSource',
'vf_Make',
'vf_MakeID',
'vf_Manufacturer',
'vf_ManufacturerId',
'vf_Model',
'vf_ModelID',
'vf_ModelYear',
'vf_MotorcycleChassisType',
'vf_MotorcycleSuspensionType',
'vf_NCSABodyType',
'vf_NCSAMake',
'vf_NCSAMapExcApprovedBy',
'vf_NCSAMapExcApprovedOn',
'vf_NCSAMappingException',
'vf_NCSAModel',
'vf_NCSANote',
'vf_Note',
'vf_OtherBusInfo',
'vf_OtherEngineInfo',
'vf_OtherMotorcycleInfo',
'vf_OtherRestraintSystemInfo',
'vf_OtherTrailerInfo',
'vf_ParkAssist',
'vf_PedestrianAutomaticEmergencyBraking',
'vf_PlantCity',
'vf_PlantCompanyName',
'vf_PlantCountry',
'vf_PlantState',
'vf_PossibleValues',
'vf_Pretensioner',
'vf_RearCrossTrafficAlert',
'vf_RearVisibilitySystem',
'vf_SAEAutomationLevel',
'vf_SAEAutomationLevel_to',
'vf_SeatBeltsAll',
'vf_SeatRows',
'vf_Seats',
'vf_SemiautomaticHeadlampBeamSwitching',
'vf_Series',
'vf_Series2',
'vf_SteeringLocation',
'vf_SuggestedVIN',
'vf_TPMS',
'vf_TopSpeedMPH',
'vf_TrackWidth',
'vf_TractionControl',
'vf_TrailerBodyType',
'vf_TrailerLength',
'vf_TrailerType',
'vf_TransmissionSpeeds',
'vf_TransmissionStyle',
'vf_Trim',
'vf_Trim2',
'vf_Turbo',
'vf_VIN',
'vf_ValveTrainDesign',
'vf_VehicleType',
'vf_WheelBaseLong',
```

```
'vf_WheelBaseShort',
'vf_WheelBaseType',
'vf_WheelSizeFront',
'vf_WheelSizeRear',
'vf_Wheels',
'vf_Windows']
```

```
In [8]: from pyspark.sql.types import StringType, BooleanType, IntegerType
```

```
In [9]: df_pyspark = df.select(['vin','firstseen','lastseen','askPrice','mileage','isNew','brandName','modelName',
    'vf_AirBagLocFront','vf_AirBagLocKnee','vf_AirBagLocSide','vf_Axles','vf_BasePrice','vf_Displacement',
    'vf_EngineCylinders','vf_EngineHP','vf_EngineKW','vf_SeatRows','vf_Seats','vf_TopSpeedMPH','vf_TransmissionSpeeds',
    'vf_WheelBaseShort','vf_WheelSizeFront','vf_WheelSizeRear','vf_Wheels','vf_Windows','msrp'])
```

```
In [10]: df_pyspark
```

```
Out[10]: DataFrame[vin: string, firstseen: timestamp, lastseen: timestamp, askPrice: int, mileage: int, isNew:
boolean, brandName: string, modelName: string, vf_AirBagLocFront: string, vf_AirBagLocKnee: string, vf_
_AirBagLocSide: string, vf_Axles: int, vf_BasePrice: double, vf_DisplacementCC: double, vf_DisplacementCI: double, vf_DisplacementL: double, vf_Doors: int, vf_EngineCylinders: int, vf_EngineHP: double, vf_EngineKW: double, vf_SeatRows: int, vf_Seats: int, vf_TopSpeedMPH: int, vf_TransmissionSpeeds: int, vf_WheelBaseShort: double, vf_WheelSizeFront: int, vf_WheelSizeRear: int, vf_Wheels: int, vf_Windows: int, msrp: int]
```

```
In [11]: #Q1
from pyspark.sql.functions import datediff,col,lit
df1 = df_pyspark.withColumn('Time in lot',datediff(df['lastSeen'],df['FirstSeen']))
```

```
In [12]: #dropping first seen and last seen columns as we have already added a new column 'Time in Lot' for it
df2 = df1.drop(col("firstseen"))
df2 = df2.drop(col("lastseen"))
df2.dtypes
```

```
Out[12]: [('vin', 'string'),
('askPrice', 'int'),
('mileage', 'int'),
('isNew', 'boolean'),
('brandName', 'string'),
('modelName', 'string'),
('vf_AirBagLocFront', 'string'),
('vf_AirBagLocKnee', 'string'),
('vf_AirBagLocSide', 'string'),
('vf_Axles', 'int'),
('vf_BasePrice', 'double'),
('vf_DisplacementCC', 'double'),
('vf_DisplacementCI', 'double'),
('vf_DisplacementL', 'double'),
('vf_Doors', 'int'),
('vf_EngineCylinders', 'int'),
('vf_EngineHP', 'double'),
('vf_EngineKW', 'double'),
('vf_SeatRows', 'int'),
('vf_Seats', 'int'),
('vf_TopSpeedMPH', 'int'),
('vf_TransmissionSpeeds', 'int'),
('vf_WheelBaseShort', 'double'),
('vf_WheelSizeFront', 'int'),
('vf_WheelSizeRear', 'int'),
('vf_Wheels', 'int'),
('vf_Windows', 'int'),
('msrp', 'int'),
('Time in lot', 'int')]
```

```
In [13]: from pyspark.sql.functions import mean
df2 = df2.na.drop(subset=['msrp','vf_BasePrice','askPrice'])
```

```
In [14]: df2.select('msrp','vf_BasePrice','askPrice').show()
```

```

+-----+-----+-----+
| msrp|vf_BasePrice|askPrice|
+-----+-----+-----+
|12387|    23475.0|   12387|
|12970|    26500.0|   12970|
|15218|    23940.0|   15218|
|18755|    38495.0|   18755|
|36999|    34175.0|   36999|
|18276|    24105.0|   18276|
|22140|    28000.0|   22140|
|    0|    34175.0|    0|
|20494|    30530.0|   20494|
|19026|    26500.0|   19026|
|18951|    26500.0|   18951|
|16649|    27995.0|   16649|
|16671|    30420.0|   16671|
|16998|    27995.0|   16998|
|20294|    30530.0|   20294|
|45996|    63000.0|   45996|
|23537|    29995.0|   23537|
|12999|    23740.0|   12999|
|13398|    23740.0|   13398|
|14121|    19995.0|   13692|
+-----+-----+-----+
only showing top 20 rows

```

```

In [15]: import pyspark.sql.functions as F
df3 = df2.withColumn('isNew', F.when(df2['isNew'] == 'FALSE', 0).otherwise(1))
df3.select('isNew').show()

```

```

+-----+
|isNew|
+-----+
|    0|
|    0|
|    0|
|    0|
|    0|
|    1|
|    1|
|    1|
|    0|
|    1|
|    1|
|    0|
|    0|
|    0|
|    0|
|    0|
|    1|
|    0|
|    0|
|    0|
+-----+
only showing top 20 rows

```

```

In [16]: #changing vf_AirBagLocFront from string to integer
from pyspark.sql.functions import when
df4 = df3.withColumn('vf_AirBagLocFront', (when(df3['vf_AirBagLocFront'] == '1st Row (Driver & Passenger)', 1).otherwise(0)))
df4.select('vf_AirBagLocFront').show(45)

```

only showing top 45 rows

```
In [17]: #changing vf_AirBagLocSide from string to integer
df5 = df4.withColumn('vf_AirBagLocSide', (when(df4['vf_AirBagLocSide'] == '1st Row (Driver & Passenger)'
df5.select('vf_AirBagLocSide').show()
```

+-----+	
vf_AirBagLocSide	
+-----+	
	4
	4
	4
	6
	6
	4
	4
	6
	6
	4
	4
	4
	4
	6
	2
	4
	4
	4
	4
+-----+	

only showing top 20 rows

```
In [18]: #changing vf_AirBagLocKnee from string to integer
df6 = df5.withColumn('vf_AirBagLocKnee', (when(df5['vf_AirBagLocKnee'] == '1st Row (Driver & Passenger)'
df6.select('vf_AirBagLocKnee').show()
```

+-----+	
vf_AirBagLocKnee	
+-----+	
	2
	2
	0
	0
	0
	0
	0
	0
	0
	0
	0
	0
	2
	0
	0
	0
	2
	2
	0
+-----+	

only showing top 20 rows

```
In [19]: df6.select('vf_Axles').show()
```

```

+-----+
|vf_Axles|
+-----+
|      2|
|      2|
|      2|
|      2|
|      2|
|      2|
|      2|
|      2|
|      2|
|      2|
|      2|
|      2|
|      2|
|      2|
|      2|
|      2|
|      2|
|      2|
|      2|
+-----+

```

only showing top 20 rows

```

In [20]: #vf_Axles - filling null values with average
avg_axles = df6.agg({'vf_Axles': 'mean' })
df7 = df6.na.fill(value=avg_axles.first()[0],subset=["vf_Axles"])
df7.select('vf_Axles').show()

```

```

+-----+
|vf_Axles|
+-----+
|      2|
|      2|
|      2|
|      2|
|      2|
|      2|
|      2|
|      2|
|      2|
|      2|
|      2|
|      2|
|      2|
|      2|
|      2|
|      2|
|      2|
|      2|
|      2|
+-----+

```

only showing top 20 rows

```

In [21]: df7.select('vf_BasePrice').show()

```



```

+-----+
|vf_BasePrice|
+-----+
|    23475.0|
|    26500.0|
|    23940.0|
|    38495.0|
|    34175.0|
|    24105.0|
|    28000.0|
|    34175.0|
|    30530.0|
|    26500.0|
|    26500.0|
|    27995.0|
|    30420.0|
|    27995.0|
|    30530.0|
|    63000.0|
|    29995.0|
|    23740.0|
|    23740.0|
|    19995.0|
+-----+

```

only showing top 20 rows

```

In [22]: #vf_BasePrice - filling null values with average
avg_bprice = df7.agg({'vf_BasePrice': 'mean' })
df8 = df7.na.fill(value=int(avg_bprice.first()[0]),subset=["vf_BasePrice"])
df8.select('vf_BasePrice').show()

```

```

+-----+
|vf_BasePrice|
+-----+
|    23475.0|
|    26500.0|
|    23940.0|
|    38495.0|
|    34175.0|
|    24105.0|
|    28000.0|
|    34175.0|
|    30530.0|
|    26500.0|
|    26500.0|
|    27995.0|
|    30420.0|
|    27995.0|
|    30530.0|
|    63000.0|
|    29995.0|
|    23740.0|
|    23740.0|
|    19995.0|
+-----+

```

only showing top 20 rows

```

In [23]: #vf_SeatRows, vf_Seats, vf_TopSpeedMPH, vf_TransmissionSpeeds, vf_WheelBaseShort, vf_WheelSizeFront,
#vf_WheelSizeRear, vf_Wheels, vf_Windows - filling null values with average
avg_dcc = df8.agg({'vf_DisplacementCC': 'mean' })
df9 = df8.na.fill(value=int(avg_dcc.first()[0]),subset=["vf_DisplacementCC"])
avg_dci = df9.agg({'vf_DisplacementCI': 'mean' })
df9 = df9.na.fill(value=int(avg_dci.first()[0]),subset=["vf_DisplacementCI"])
avg_dl = df9.agg({'vf_DisplacementL': 'mean' })
df9 = df9.na.fill(value=int(avg_dl.first()[0]),subset=["vf_DisplacementL"])
df9 = df9.na.fill(value=4,subset=["vf_Doors"])
avg_cyl = df9.agg({'vf_EngineCylinders': 'mean' })
df9 = df9.na.fill(value=int(avg_cyl.first()[0]),subset=["vf_EngineCylinders"])
avg_ehp = df9.agg({'vf_EngineHP': 'mean' })
df9 = df9.na.fill(value=avg_ehp.first()[0],subset=["vf_EngineHP"])
avg_ekw = df9.agg({'vf_EngineKW': 'mean' })
df9 = df9.na.fill(value=avg_ekw.first()[0],subset=["vf_EngineKW"])
avg_sr = df9.agg({'vf_SeatRows': 'mean' })

```

```
df9 = df9.na.fill(value=int(avg_sr.first()[0]),subset=["vf_SeatRows"])
avg_seats = df9.agg({'vf_Seats': 'mean' })
df9 = df9.na.fill(value=int(avg_seats.first()[0]),subset=["vf_Seats"])
avg_speed = df9.agg({'vf_TopSpeedMPH': 'mean' })
df9 = df9.na.fill(value=int(avg_speed.first()[0]),subset=["vf_TopSpeedMPH"])
avg_tspeed = df9.agg({'vf_TransmissionSpeeds': 'mean' })
df9 = df9.na.fill(value=int(avg_tspeed.first()[0]),subset=["vf_TransmissionSpeeds"])
avg_wbs = df9.agg({'vf_WheelBaseShort': 'mean' })
df9 = df9.na.fill(value=avg_wbs.first()[0],subset=["vf_WheelBaseShort"])
avg_wsf = df9.agg({'vf_WheelSizeFront': 'mean' })
df9 = df9.na.fill(value=int(avg_wsf.first()[0]),subset=["vf_WheelSizeFront"])
avg_wsr = df9.agg({'vf_WheelSizeRear': 'mean' })
df9 = df9.na.fill(value=int(avg_wsr.first()[0]),subset=["vf_WheelSizeRear"])
avg_wheels = df9.agg({'vf_Wheels': 'mean' })
df9 = df9.na.fill(value=int(avg_wheels.first()[0]),subset=["vf_Wheels"])
df9 = df9.na.fill(value=4,subset=["vf_Windows"])
df9.select('vf_DisplacementCC','vf_DisplacementCI','vf_DisplacementL','vf_Doors','vf_EngineCylinders',
```


In [24]: `df9.dtypes`

Out[24]:

```
(('vin', 'string'),
 ('askPrice', 'int'),
 ('mileage', 'int'),
 ('isNew', 'int'),
 ('brandName', 'string'),
 ('modelName', 'string'),
 ('vf_AirBagLocFront', 'int'),
 ('vf_AirBagLocKnee', 'int'),
 ('vf_AirBagLocSide', 'int'),
 ('vf_Axles', 'int'),
 ('vf_BasePrice', 'double'),
 ('vf_DisplacementCC', 'double'),
 ('vf_DisplacementCI', 'double'),
 ('vf_DisplacementL', 'double'),
 ('vf_Doors', 'int'),
 ('vf_EngineCylinders', 'int'),
 ('vf_EngineHP', 'double'),
 ('vf_EngineKW', 'double'),
 ('vf_SeatRows', 'int'),
 ('vf_Seats', 'int'),
 ('vf_TopSpeedMPH', 'int'),
 ('vf_TransmissionSpeeds', 'int'),
 ('vf_WheelBaseShort', 'double'),
 ('vf_WheelSizeFront', 'int'),
 ('vf_WheelSizeRear', 'int'),
 ('vf_Wheels', 'int'),
 ('vf_Windows', 'int'),
 ('msrp', 'int'),
 ('Time in lot', 'int'))
```

In [25]:

```
from pyspark.ml.feature import VectorAssembler
ip_cols = ['askPrice', 'mileage', 'isNew', 'vf_AirBagLocFront', 'vf_AirBagLocKnee', 'vf_AirBagLocSide', 'vf_
vf_EngineCylinders', 'vf_EngineHP', 'vf_EngineKW', 'vf_SeatRows', 'vf_Seats', 'vf_TopSpeedMPH', 'vf_Trans
op_col = "Features"
vec_df = VectorAssembler(inputCols = ip_cols, outputCol = op_col)

df_ml = vec_df.transform(df9)
df_ml.select(['Features']).toPandas().head(5)
```

Out[25]:

	Features
0	[12387.0, 0.0, 0.0, 2.0, 2.0, 4.0, 2.0, 23475....
1	[12970.0, 0.0, 0.0, 2.0, 2.0, 4.0, 2.0, 26500....
2	[15218.0, 0.0, 0.0, 2.0, 0.0, 4.0, 2.0, 23940....
3	[18755.0, 0.0, 0.0, 2.0, 0.0, 6.0, 2.0, 38495....
4	[36999.0, 0.0, 0.0, 2.0, 0.0, 6.0, 2.0, 34175....

In [26]:

```
final_df = df_ml.select(['Features', 'msrp'])
final_df.show(1)
```

```
+-----+-----+-----+
|           Features| msrp|
+-----+-----+-----+
|[12387.0,0.0,0.0,...]|12387|
+-----+-----+-----+
only showing top 1 row
```

In [27]:

```
len_df = final_df.count()

train_len = int(0.7*len_df)
train_len
```

Out[27]: 1299720

In [28]:

```
test_df = final_df
train_df = final_df.limit(train_len)
train_df.show()
train_df.count()
```

```

+-----+-----+
|          Features| msrp|
+-----+-----+
|[12387.0,0.0,0.0,...|12387|
|[12970.0,0.0,0.0,...|12970|
|[15218.0,0.0,0.0,...|15218|
|[18755.0,0.0,0.0,...|18755|
|[36999.0,0.0,0.0,...|36999|
|[18276.0,0.0,1.0,...|18276|
|[22140.0,0.0,1.0,...|22140|
|[0.0,0.0,1.0,2.0,...| 0|
|[20494.0,0.0,0.0,...|20494|
|[19026.0,0.0,1.0,...|19026|
|[18951.0,0.0,1.0,...|18951|
|[16649.0,0.0,0.0,...|16649|
|[16671.0,0.0,0.0,...|16671|
|[16998.0,0.0,0.0,...|16998|
|[20294.0,0.0,0.0,...|20294|
|[45996.0,0.0,0.0,...|45996|
|[23537.0,0.0,1.0,...|23537|
|[12999.0,0.0,0.0,...|12999|
|[13398.0,0.0,0.0,...|13398|
|[13692.0,0.0,0.0,...|14121|
+-----+-----+
only showing top 20 rows

```

Out[28]: 1299720

```

In [29]: test_df = test_df.subtract(train_df)
test_df.show()
test_df.count()

```

```

+-----+-----+
|          Features| msrp|
+-----+-----+
|[0.0,0.0,1.0,2.0,...| 0|
|[0.0,0.0,1.0,2.0,...| 0|
|[0.0,0.0,1.0,2.0,...| 0|
|[0.0,0.0,1.0,2.0,...| 0|
|[0.0,0.0,1.0,2.0,...| 0|
|[0.0,3.0,1.0,2.0,...| 0|
|[0.0,9.0,1.0,2.0,...| 0|
|[0.0,12153.0,0.0,...| 0|
|[0.0,15144.0,0.0,...| 0|
|[11499.0,45603.0,...|12399|
|[12928.0,0.0,1.0,...|16828|
|[12995.0,0.0,0.0,...|14031|
|[13500.0,0.0,0.0,...|14500|
|[14309.0,27687.0,...|15000|
|[14950.0,9905.0,0...|14950|
|[14994.0,0.0,0.0,...|16975|
|[15303.0,38997.0,...|15599|
|[15395.0,0.0,0.0,...|16995|
|[15626.0,0.0,0.0,...|15726|
|[15919.0,0.0,1.0,...|19570|
+-----+-----+
only showing top 20 rows

```

Out[29]: 360992

```

In [30]: train_df.count()

```

Out[30]: 1299720

```

In [31]: test_df.count()

```

Out[31]: 360992

```

In [32]: from pyspark.ml.regression import LinearRegression

```

```

In [33]: lr = LinearRegression(featuresCol = 'Features', labelCol='msrp', maxIter=1000)

```

```
In [34]: train_df.show(100)
```

+-----+-----+	
Features	msrp
+-----+-----+	
[12387.0,0.0,0.0,...]	12387
[12970.0,0.0,0.0,...]	12970
[15218.0,0.0,0.0,...]	15218
[18755.0,0.0,0.0,...]	18755
[36999.0,0.0,0.0,...]	36999
[18276.0,0.0,1.0,...]	18276
[22140.0,0.0,1.0,...]	22140
[0.0,0.0,1.0,2.0,...]	0
[20494.0,0.0,0.0,...]	20494
[19026.0,0.0,1.0,...]	19026
[18951.0,0.0,1.0,...]	18951
[16649.0,0.0,0.0,...]	16649
[16671.0,0.0,0.0,...]	16671
[16998.0,0.0,0.0,...]	16998
[20294.0,0.0,0.0,...]	20294
[45996.0,0.0,0.0,...]	45996
[23537.0,0.0,1.0,...]	23537
[12999.0,0.0,0.0,...]	12999
[13398.0,0.0,0.0,...]	13398
[13692.0,0.0,0.0,...]	14121
[20997.0,0.0,0.0,...]	20997
[12568.0,0.0,1.0,...]	12568
[14680.0,0.0,1.0,...]	15680
[16829.0,0.0,1.0,...]	17829
[12568.0,0.0,1.0,...]	12568
[12959.0,0.0,0.0,...]	12959
[12998.0,0.0,0.0,...]	12998
[19514.0,0.0,1.0,...]	19514
[12930.0,0.0,0.0,...]	12930
[12970.0,0.0,0.0,...]	12994
[20978.0,0.0,0.0,...]	20978
[23652.0,0.0,1.0,...]	23652
[10999.0,0.0,0.0,...]	10999
[12575.0,0.0,0.0,...]	12575
[24457.0,0.0,0.0,...]	24457
[12799.0,0.0,0.0,...]	12799
[14529.0,0.0,0.0,...]	14529
[15683.0,0.0,0.0,...]	15683
[24726.0,0.0,0.0,...]	24726
[12572.0,0.0,1.0,...]	12572
[17999.0,0.0,0.0,...]	17999
[45644.0,0.0,0.0,...]	45644
[12568.0,0.0,1.0,...]	12568
[18868.0,0.0,1.0,...]	19868
[24286.0,0.0,1.0,...]	24286
[12394.0,0.0,0.0,...]	12394
[14520.0,0.0,0.0,...]	14999
[24885.0,0.0,0.0,...]	24885
[24998.0,0.0,0.0,...]	24998
[11999.0,0.0,0.0,...]	11999
[15976.0,0.0,1.0,...]	16476
[13769.0,0.0,0.0,...]	13819
[12902.0,0.0,0.0,...]	12902
[17645.0,0.0,0.0,...]	17695
[12357.0,0.0,1.0,...]	12357
[14384.0,0.0,0.0,...]	14384
[18781.0,0.0,1.0,...]	18781
[23550.0,0.0,1.0,...]	23550
[13295.0,0.0,0.0,...]	13295
[18328.0,0.0,0.0,...]	18328
[26060.0,0.0,0.0,...]	26060
[12949.0,0.0,0.0,...]	12999
[16925.0,0.0,1.0,...]	16925
[15900.0,0.0,0.0,...]	15986
[12899.0,0.0,0.0,...]	12900
[17821.0,0.0,0.0,...]	17821
[33841.0,0.0,1.0,...]	33841
[33125.0,0.0,1.0,...]	33125
[12998.0,0.0,0.0,...]	12998
[14621.0,0.0,0.0,...]	14667
[20383.0,0.0,1.0,...]	20383
[29288.0,0.0,1.0,...]	29288

```
[36163.0,0.0,1.0,...|36163|
[16860.0,0.0,1.0,...|16860|
[12994.0,0.0,0.0,...|12994|
[12995.0,0.0,0.0,...|13509|
[12281.0,0.0,0.0,...|12685|
[16462.0,0.0,1.0,...|17462|
[23481.0,0.0,1.0,...|23481|
[15999.0,0.0,0.0,...|16475|
[18787.0,0.0,1.0,...|23787|
[16648.0,0.0,0.0,...|16648|
[18787.0,0.0,1.0,...|18787|
[12687.0,0.0,0.0,...|12951|
[12998.0,0.0,0.0,...|13596|
[18851.0,0.0,0.0,...|18875|
[16860.0,0.0,1.0,...|21360|
[18787.0,0.0,1.0,...|20037|
[20663.0,0.0,1.0,...|24163|
[12407.0,0.0,1.0,...|12407|
[12568.0,0.0,1.0,...|12568|
[16860.0,0.0,1.0,...|21360|
[24999.0,0.0,0.0,...|24999|
[13998.0,0.0,0.0,...|13998|
[21689.0,0.0,0.0,...|21689|
[14800.0,0.0,0.0,...|14800|
[16349.0,0.0,0.0,...|16428|
[18272.0,0.0,0.0,...|18273|
[10339.0,0.0,0.0,...|10752|
[21935.0,0.0,0.0,...|21935|
+-----+
only showing top 100 rows
```

In [35]: `test_df.show()`

```
+-----+-----+-----+
|          Features| msrp|
+-----+-----+-----+
|[0.0,0.0,1.0,2.0,...| 0|
|[0.0,0.0,1.0,2.0,...| 0|
|[0.0,0.0,1.0,2.0,...| 0|
|[0.0,0.0,1.0,2.0,...| 0|
|[0.0,0.0,1.0,2.0,...| 0|
|[0.0,3.0,1.0,2.0,...| 0|
|[0.0,9.0,1.0,2.0,...| 0|
|[0.0,12153.0,0.0,...| 0|
|[0.0,15144.0,0.0,...| 0|
|[11499.0,45603.0,...|12399|
|[12928.0,0.0,1.0,...|16828|
|[12995.0,0.0,0.0,...|14031|
|[13500.0,0.0,0.0,...|14500|
|[14309.0,27687.0,...|15000|
|[14950.0,9905.0,0...|14950|
|[14994.0,0.0,0.0,...|16975|
|[15303.0,38997.0,...|15599|
|[15395.0,0.0,0.0,...|16995|
|[15626.0,0.0,0.0,...|15726|
|[15919.0,0.0,1.0,...|19570|
+-----+-----+-----+
only showing top 20 rows
```

In [36]: `lr_model = lr.fit(train_df)`

In [37]: `y_pred = lr_model.transform(test_df)`

In [38]: `y_pred.show()`

Features	msrp	prediction
[0.0,0.0,1.0,2.0,...]	0	1125.0722357587292
[0.0,0.0,1.0,2.0,...]	0	16652.786260337907
[0.0,0.0,1.0,2.0,...]	0	-1202.6032615360382
[0.0,0.0,1.0,2.0,...]	0	9599.820605161342
[0.0,0.0,1.0,2.0,...]	0	-15991.774871472644
[0.0,3.0,1.0,2.0,...]	0	-17861.4783095269
[0.0,9.0,1.0,2.0,...]	0	18427.370204653864
[0.0,12153.0,0.0,...]	0	13243.816229516342
[0.0,15144.0,0.0,...]	0	13290.599226041759
[11499.0,45603.0,...]	12399	42487.73347571959
[12928.0,0.0,1.0,...]	16828	27972.138046448712
[12995.0,0.0,0.0,...]	14031	40274.64985817963
[13500.0,0.0,0.0,...]	14500	31408.758727997727
[14309.0,27687.0,...]	15000	42971.42268637994
[14950.0,9905.0,0.0,...]	14950	21933.404842900036
[14994.0,0.0,0.0,...]	16975	43991.068482176444
[15303.0,38997.0,...]	15599	39046.84608497289
[15395.0,0.0,0.0,...]	16995	16904.26493510482
[15626.0,0.0,0.0,...]	15726	24146.79598416574
[15919.0,0.0,1.0,...]	19570	12697.472309073622

only showing top 20 rows

```
In [39]: y_pred.describe().show()
```

summary	msrp	prediction
count	360992	360992
mean	2000863.2584406303	446060.76193985925
stddev	6.499651648471686E7	3.225782477478744E7
min	0	-72687.01466037614
max	2147483647	2.5674029905667214E9

```
In [40]: print("Coefficients: " + str(lr_model.coefficients))
print("Intercept: " + str(lr_model.intercept))
```

```
Coefficients: [1.1955130707373247,0.39095921094982955,-4256.668333671143,-960.4479598289017,-4513.4733
53312087,4829.620689285545,0.0,-0.02818322086852489,-0.026681537555876916,-0.45795822902252364,-19.074
779173113654,3924.2135531015383,-2861.0278897925405,-54.87006199341723,36.44750887391078,-4903.3861310
88277,-2160.188891211929,290.9167649494918,-57.52588882733706,-0.5990587556657906,3441.0627458695367,-
2841.2497410185674,3197.880073594043,-5646.643740351718,84.31609582145248]
Intercept: -19468.11755917413
```

```
In [41]: trainingSummary = lr_model.summary
print("RMSE: %f" % trainingSummary.rootMeanSquaredError)
print("r2: %f" % trainingSummary.r2)
```

```
RMSE: 3938430.054993
r2: 0.000040
```

```
In [42]: y_pred.select("prediction","msrp","Features").show()
from pyspark.ml.evaluation import RegressionEvaluator
lr_evaluator = RegressionEvaluator(predictionCol="prediction", labelCol="msrp",metricName="r2")
print("R Squared (R2) on test data = %g" % lr_evaluator.evaluate(y_pred))
```

prediction	msrp	Features
1125.0722357587292	0	[0.0,0.0,1.0,2.0,...]
16652.786260337907	0	[0.0,0.0,1.0,2.0,...]
-1202.6032615360382	0	[0.0,0.0,1.0,2.0,...]
9599.820605161342	0	[0.0,0.0,1.0,2.0,...]
-15991.774871472644	0	[0.0,0.0,1.0,2.0,...]
-17861.4783095269	0	[0.0,3.0,1.0,2.0,...]
18427.370204653864	0	[0.0,9.0,1.0,2.0,...]
13243.816229516342	0	[0.0,12153.0,0.0,...]
13290.599226041759	0	[0.0,15144.0,0.0,...]
42487.73347571959	12399	[11499.0,45603.0,...]
27972.138046448712	16828	[12928.0,0.0,1.0,...]
40274.64985817963	14031	[12995.0,0.0,0.0,...]
31408.758727997727	14500	[13500.0,0.0,0.0,...]
42971.42268637994	15000	[14309.0,27687.0,...]
21933.404842900036	14950	[14950.0,9905.0,0...]
43991.068482176444	16975	[14994.0,0.0,0.0,...]
39046.84608497289	15599	[15303.0,38997.0,...]
16904.26493510482	16995	[15395.0,0.0,0.0,...]
24146.79598416574	15726	[15626.0,0.0,0.0,...]
12697.472309073622	19570	[15919.0,0.0,1.0,...]

only showing top 20 rows

R Squared (R2) on test data = 0.164857

```
In [43]: #Decision tree
from pyspark.ml.regression import DecisionTreeRegressor
dt = DecisionTreeRegressor(featuresCol='Features', labelCol='msrp')
dt_model = dt.fit(train_df)
dt_pred = dt_model.transform(test_df)
dt_evaluator = RegressionEvaluator(
    labelCol="msrp", predictionCol="prediction", metricName="rmse")
rmse = dt_evaluator.evaluate(dt_pred)
dt_pred.select("prediction", "msrp", "Features").show(5)
print("Root Mean Squared Error (RMSE) on test data = %g" % rmse)
```

prediction	msrp	Features
0.0	0	[0.0,0.0,1.0,2.0,...]
0.0	0	[0.0,0.0,1.0,2.0,...]
0.0	0	[0.0,0.0,1.0,2.0,...]
0.0	0	[0.0,0.0,1.0,2.0,...]
0.0	0	[0.0,0.0,1.0,2.0,...]

only showing top 5 rows

Root Mean Squared Error (RMSE) on test data = 6.5025e+07

```
In [44]: #Gradient boosted tree regression
from pyspark.ml.regression import GBRegressor
gbt = GBRegressor(featuresCol='Features', labelCol='msrp', maxIter=10)
gbt_model = gbt.fit(train_df)
gbt_predictions = gbt_model.transform(test_df)
gbt_predictions.select('prediction', 'msrp', 'Features').show(5)
```

prediction	msrp	Features
1049.659896873842	0	[0.0,0.0,1.0,2.0,...]
-790.2891629690153	0	[0.0,0.0,1.0,2.0,...]
232.09748683109802	0	[0.0,0.0,1.0,2.0,...]
1124.774128229476	0	[0.0,0.0,1.0,2.0,...]
95.70418929729463	0	[0.0,0.0,1.0,2.0,...]

only showing top 5 rows

```
In [45]: gbt_evaluator = RegressionEvaluator(
    labelCol="msrp", predictionCol="prediction", metricName="rmse")
rmse = gbt_evaluator.evaluate(gbt_predictions)
print("Root Mean Squared Error (RMSE) on test data = %g" % rmse)
```

Root Mean Squared Error (RMSE) on test data = 6.50247e+07

```
In [46]: #Random Forest regression

from pyspark.ml.regression import RandomForestRegressor
rf = RandomForestRegressor(featuresCol = 'Features', labelCol = 'msrp')
rf_model = rf.fit(train_df)
rf_predictions = rf_model.transform(test_df)
rf_predictions.select('prediction', 'msrp', 'Features').show(5)
```

```
+-----+-----+
| prediction|msrp| Features|
+-----+-----+
|10450.204172665788| 0|[0.0,0.0,1.0,2.0,...|
|12339.358591085871| 0|[0.0,0.0,1.0,2.0,...|
|10912.716874344278| 0|[0.0,0.0,1.0,2.0,...|
| 17527.58963664177| 0|[0.0,0.0,1.0,2.0,...|
| 17342.74259213676| 0|[0.0,0.0,1.0,2.0,...|
+-----+-----+
only showing top 5 rows
```

```
In [47]: rf_evaluator = RegressionEvaluator(
    labelCol="msrp", predictionCol="prediction", metricName="rmse")
rmse = rf_evaluator.evaluate(rf_predictions)
print("Root Mean Squared Error (RMSE) on test data = %g" % rmse)
```

Root Mean Squared Error (RMSE) on test data = 6.50259e+07

```
In [48]: #isotonic regressor

from pyspark.ml.regression import IsotonicRegression
ir = IsotonicRegression(featuresCol = 'Features', labelCol = 'msrp')
ir_model = ir.fit(train_df)
ir_predictions = ir_model.transform(test_df)
ir_predictions.select('prediction', 'msrp', 'Features').show(5)
```

```
+-----+-----+
| prediction|msrp| Features|
+-----+-----+
| 0.0| 0|[0.0,0.0,1.0,2.0,...|
| 0.0| 0|[0.0,0.0,1.0,2.0,...|
| 0.0| 0|[0.0,0.0,1.0,2.0,...|
| 0.0| 0|[0.0,0.0,1.0,2.0,...|
| 0.0| 0|[0.0,0.0,1.0,2.0,...|
+-----+-----+
only showing top 5 rows
```

```
In [49]: ir_evaluator = RegressionEvaluator(
    labelCol="msrp", predictionCol="prediction", metricName="rmse")
rmse = ir_evaluator.evaluate(ir_predictions)
print("Root Mean Squared Error (RMSE) on test data = %g" % rmse)
```

Root Mean Squared Error (RMSE) on test data = 6.49957e+07

Specific Question Proposed

```
In [50]: #
df = df.select('firstSeen',
    'lastSeen',
    'msrp',
    'askPrice',
    'mileage',
    (df.isNew.cast(IntegerType())),      ## Converting Boolean to Numeric
    'color',
    'brandName',
    'modelName')
```

```
In [51]: df.count()
```

```
Out[51]: 5695015
```

```
In [52]: df.na.drop(how = "any")
```

```
Out[52]: DataFrame[firstSeen: timestamp, lastSeen: timestamp, msrp: int, askPrice: int, mileage: int, isNew: in
t, color: string, brandName: string, modelName: string]
```

```
In [53]: df_car_pyspark = df.sample(0.1)
```

```
In [54]: df_car_pyspark.count()
```

```
Out[54]: 570977
```

```
In [55]: df_car_pyspark.printSchema()
```

```
root
 |-- firstSeen: timestamp (nullable = true)
 |-- lastSeen: timestamp (nullable = true)
 |-- msrp: integer (nullable = true)
 |-- askPrice: integer (nullable = true)
 |-- mileage: integer (nullable = true)
 |-- isNew: integer (nullable = true)
 |-- color: string (nullable = true)
 |-- brandName: string (nullable = true)
 |-- modelName: string (nullable = true)
```

```
In [56]: df_car_pyspark.describe().show()
```

```
+-----+-----+-----+-----+-----+
|summary|msrp|askPrice|mileage|isNew|
color|brandName|modelName|
+-----+-----+-----+-----+-----+
|count|570977|570977|570977|570977|5
70977|570861|570371|
|mean|720776.410953506|177120.53963119356|22006.17455869501|0.34448497925485616|16413.8904109
58906|null|1352.6327305489415|
|stddev|3.864848769443795E7|1.819671471510693E7|45503.31890962189|0.47520045645206654|49950.70672
27175|null|835.4543171304593|
|min|0|0|0|0|
|ACURA|124 Spider|
|max|2147483647|2147483647|9024018|1|white silver me
ta...|YAMAHA|new Passat|
+-----+-----+-----+-----+-----+
```

```
In [57]: #Q1 How long brfore the car is sold?
```

```
from pyspark.sql.functions import datediff,col,lit
```

```
df1 = df_car_pyspark.withColumn('Time_in_lot',datediff(df_car_pyspark['lastSeen'],df_car_pyspark['First
```

```
In [58]: df1 = df1.select('Time_in_lot')
```

```
In [59]: df1.dtypes
```

```
Out[59]: [('Time_in_lot', 'int')]
```

```
In [60]: df1.sort(df1.Time_in_lot.desc()).show(5)
```

```
+-----+
|Time_in_lot|
+-----+
|1396|
|1395|
|1369|
|1365|
|1364|
+-----+
only showing top 5 rows
```

```
In [61]: df1.describe().show()
```

```

+-----+-----+
|summary|      Time_in_lot|
+-----+-----+
|  count|          570977|
|   mean| 67.92468172973692|
|  stddev|106.31335378028105|
|    min|          -1032|
|    max|           1396|
+-----+-----+

```

As seen above, a summary of the timeInLot variable (measured in days) shows that the minimum number of days that a car was in a dealer's lot was zero, whereas the max was 1,409 days (~3.9 years).

In [62]: *#Q2 Which Cars are the Most Popular?*

```
from pyspark.sql.functions import concat,col,lit,asc,desc
```

In [63]: `df2 = df_car_pyspark.select(concat(col("brandName"),lit(" "),col("modelName")).alias("Cars"))`

In [64]: `df2.show(10)`

```

+-----+-----+
|          Cars|
+-----+-----+
|MITSUBISHI Eclips...|
|      NISSAN Altima|
|    FORD Explorer|
|      FORD Edge|
|    FORD Fusion|
|    FORD Fusion|
|    FORD Fusion|
|    FORD Fiesta|
|    LINCOLN MKX|
|    FORD Fiesta|
+-----+-----+
only showing top 10 rows

```

In [65]: `count = df2.groupBy('Cars').count()`

In [66]: `sparkcount_df = count.orderBy(col("count").desc())`

In [67]: `sparkcount_df.show(5)`

```

+-----+-----+-----+
|          Cars|count|
+-----+-----+-----+
|    FORD F-150|17366|
|CHEVROLET Silverado|16400|
|CHEVROLET Equinox|15499|
|    FORD Escape|11810|
|JEEP Grand Cherokee|10780|
+-----+-----+-----+
only showing top 5 rows

```

In [68]: `import matplotlib.pyplot as plt`
`import pandas as pd`

```
pandascount_df = sparkcount_df.toPandas()
```

In [69]: `pandascount_df.set_index('Cars',inplace =True)`


In [70]: `pandascount_df = pandascount_df.head(5)`

In [71]: `pandascount_df`

Out [71]:

	count
Cars	
FORD F-150	17366
CHEVROLET Silverado	16400
CHEVROLET Equinox	15499
FORD Escape	11810
JEEP Grand Cherokee	10780

```
In [72]: plot = pandascount_df.plot.pie(y= 'count', figsize = (10,15))
```



```
In [73]: #Q3 What Color of the car do people like the most?
```

```
df3 = df_car_pyspark.select('color')
```

```
In [74]: color = df3.groupby("color").count()
```

```
In [75]: sparkcolor_df = color.orderBy(col("count").desc())
```

```
In [76]: sparkcolor_df.show(6)
```

```
+-----+-----+
|      color| count|
+-----+-----+
|      N/A| 211012|
|     Black|  21084|
|     White| 11049|
|     Gray|   7887|
|Summit White|  7221|
|     Silver|  6794|
+-----+-----+
only showing top 6 rows
```

```
In [77]: sparkcolor_df = sparkcolor_df.filter(sparkcolor_df.color!= 'N/A')
```

```
In [78]: colorpandas_df = sparkcolor_df.toPandas()
```

```
In [79]: colorpandas_df.set_index('color',inplace =True)
```


```
In [80]: colorpandas_df = colorpandas_df.head(6)
```

```
In [81]: colorpandas_df
```

```
Out[81]:
```

	count
color	
Black	21084
White	11049
Gray	7887
Summit White	7221
Silver	6794
Blue	6171

```
In [82]: plot = colorpandas_df.plot.pie(y= 'count', figsize = (10,15))
```



In [83]: *#Q5 How much will the final price differ from MSRP on avg*

```
df5 = df_car_pyspark.select('msrp','askPrice')
```

In [84]: `df5.printSchema()`

```
root
|-- msrp: integer (nullable = true)
|-- askPrice: integer (nullable = true)
```

In [85]: `df5.show()`


```

+-----+-----+
| msrp|askPrice|
+-----+-----+
| 1498|    1498|
|10589|   10589|
|36999|   36999|
|23652|   23652|
|12575|   12575|
|15683|   15683|
|11763|   11763|
|12572|   12572|
|24885|   24885|
| 7685|    7650|
|18781|   18781|
|26060|   26060|
|36163|   36163|
|13509|   12995|
|12685|   12281|
|17321|   16993|
|20037|   18787|
|23888|   22396|
|12689|   12689|
|    0|        0|
+-----+-----+
only showing top 20 rows

```

```
In [86]: df5.corr("msrp", "askPrice")
```


```
Out[86]: 0.4707036907884924
```

```
In [87]: import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [88]: fig = plt.figure(figsize=(25, 15))
st = fig.suptitle("Distribution of Features", fontsize=50, verticalalignment="baseline")
for col, num in zip(df5.toPandas().describe().columns, range(1,11)):
    ax = fig.add_subplot(3,4, num)
    ax.hist(df5.toPandas()[col])
    plt.grid(False)
    plt.xticks(rotation=45, fontsize=20)
    plt.yticks(fontsize=15)
    plt.title(col.upper(), fontsize=20)

plt.tight_layout()
st.set_y(0.95)
fig.subplots_adjust(top=0.85, hspace=0.4)
plt.show()
```

Distribution of Features



```
In [89]: df5 = df5.withColumn('Price_diff', df5['msrp'] - df5['askPrice'])
```

```
In [90]: df5.show()
```

msrp	askPrice	Price_diff
1498	1498	0
10589	10589	0
36999	36999	0
23652	23652	0
12575	12575	0
15683	15683	0
11763	11763	0
12572	12572	0
24885	24885	0
7685	7650	35
18781	18781	0
26060	26060	0
36163	36163	0
13509	12995	514
12685	12281	404
17321	16993	328
20037	18787	1250
23888	22396	1492
12689	12689	0
0	0	0

only showing top 20 rows

```
In [91]: df5.select('price_diff').describe().show()
```

summary	price_diff
count	570977
mean	543655.8713223125
stddev	3.409922691264347E7
min	0
max	2147480657

```
In [ ]:
In [ ]:
In [ ]:
In [ ]:
In [ ]:
In [ ]:
```